

Civic Evidence and Verification

Written Summary · Track: Transparency & Accountability · Country: Kenya

The Problem

Civic information is often fragmented across official announcements, public documents, community conversations, and individual reports. Even when information is available, citizens may struggle to determine what was officially stated, what is being observed on the ground, when information was recorded, and where evidence conflicts. The problem is not only access to information. It is the gap between information and evidence.

The primary user is a resident checking whether a promised public project or service is working in their community. Journalists and watchdog groups are secondary users who need a clear, inspectable record they can investigate and discuss.

The Solution

Civic Evidence and Verification is a working Django proof of concept that creates a structured evidence layer around civic projects and public services. The platform keeps three things distinct: the official record (what an institution has stated, including its status, date, source, location, and responsible organization), the community observation (what a person reports experiencing, including the observation date, status, description, and reporter name or anonymous submission), and the supporting evidence (additional context associated with an observation, such as a photo, video, document, or testimony; the prototype records the evidence type and note but does not currently upload files).

The platform does not decide what citizens should believe. It makes the claim, observation, evidence, and disagreement visible.

How It Works

The workflow is straightforward: find a civic project or public service by name or location, review the official record and its source, compare recent community observations, and see an inspectable verdict of aligned, conflicting, stale, or no evidence yet. Contributing an observation updates the record for the next person who checks it.

The evidence logic is explicit and inspectable. It compares recent observation statuses with what the official status implies, and every verdict provides a next step: a conflicting verdict invites another observation to help clarify the disagreement, a stale verdict indicates the available observations are too old and need to be updated, an aligned verdict indicates that recent observations support the official record, and a no-evidence-yet verdict indicates there are currently no recent community observations.

Recent verdicts also show the number of distinct reporters who contributed observations. This is a transparency signal, not proof of identity or accuracy. It helps prevent a single report from being presented as equivalent to several independent reports.

Trust, Accuracy, and Information Sources

Official records and community observations remain separate. Official records show their stated source, and observations are dated and can include supporting evidence notes. Conflicting observations remain visible rather than being silently resolved, and the system distinguishes evidence from conclusions.

The seeded records are demonstration data and are clearly labelled as such throughout the interface. They are designed to demonstrate aligned, conflicting, stale, and no-evidence outcomes rather than to claim comprehensive or verified coverage of real government data.

Design for Real Conditions

Multilingual access: the interface supports English and Kiswahili through Django translation files and a language switcher. Project facts remain in their source language rather than being silently altered.

Low bandwidth: the interface uses server-rendered HTML, with no frontend framework, no large media, and no required JavaScript.

Accessibility: pages use semantic headings, labelled controls, visible keyboard focus, a skip link, and text labels alongside status colors.

Privacy: no account is required. A person can leave their name blank, and the prototype does not collect device or location data.

Local relevance: demonstration records use civic projects and public services in Kenya, with local locations and institutions.

Scalability

The data model is geography-agnostic by design. A project and an observation are generic records (status, date, location, organization, evidence type) with no Kenya-specific fields, so a new country or region is a matter of adding data, not rewriting the schema.

The interface already supports more than one language through Django's translation system, so adding a language is a translation file, not new code. Demonstration data is seeded from structured JSON files through a single management command, so a new set of civic projects for another country can be loaded the same way the Kenyan demonstration data was, without touching the application logic. The evidence-status rule is also independent of any specific project type: it compares an official status against an observation status using a fixed mapping, so new categories of civic services can be added by extending that mapping rather than rebuilding the verdict logic.

Implementation

The prototype uses Django, server-rendered templates, SQLite for local development, and a management command for demonstration data.

Use of AI Tools and Originality

The evidence-layer concept, the three-way data model, and the verification logic were specified by the author before any code was written. AI-assisted development tools built each feature under that direction: the data model, the evidence-status logic, the seed data, the translation setup, the templates, and the documentation were each specified, generated with AI assistance, and then manually reviewed and tested end to end before being accepted.

AI tools accelerated implementation and repetitive work rather than making design or product decisions. The proof of concept does not rely on an LLM as a source of truth and does not delegate the final verification decision to an AI model; that logic is implemented explicitly within the application and is fully inspectable in the code.

Why It Matters

Civic Evidence and Verification offers a practical way to make public information easier to understand, trace, and scrutinize. Its principle is simple: show the claim, show the observation, show the evidence, and make uncertainty visible.

Future Development

A larger implementation could add verified public data integrations, trusted-organization signals, geographic analysis, document ingestion, file-based evidence with appropriate safeguards, and additional civic services. These additions would build on the working evidence model and user workflow demonstrated by the current prototype.